

```
In [154]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import datetime as dt

from sklearn.model_selection import (train_test_split,
                                     cross_val_score,
                                     GridSearchCV)

from sklearn.metrics import (accuracy_score,
                              f1_score, precision_score,
                              confusion_matrix,
                              classification_report,
                              confusion_matrix, roc_curve, auc)

from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import OneHotEncoder
from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb
from sklearn.tree import DecisionTreeClassifier
```

Business Problem

Tanzania has had a problem with available water to the general populace for many years.

The Tanzanian government has hired us to figure out a way to improve methods in identifying non-functioning water wells.

We will be trying to detect which key features will help up identify the status of these wells.

Column Descriptions

amount_tsh - Total static head (amount water available to waterpoint)

date_recorded - The date the row was entered

funder - Who funded the well

gps_height - Altitude of the well

installer - Organization that installed the well

longitude - GPS coordinate

latitude - GPS coordinate

wpt_name - Name of the waterpoint if there is one

num_private -

basin - Geographic water basin

subvillage - Geographic location

region - Geographic location

region_code - Geographic location (coded)

district_code - Geographic location (coded)

lga - Geographic location

ward - Geographic location

population - Population around the well

public_meeting - True/False

recorded_by - Group entering this row of data

scheme_management - Who operates the waterpoint

scheme_name - Who operates the waterpoint

permit - If the waterpoint is permitted

construction_year - Year the waterpoint was constructed

extraction_type - The kind of extraction the waterpoint uses

extraction_type_group - The kind of extraction the waterpoint uses

extraction_type_class - The kind of extraction the waterpoint uses

management - How the waterpoint is managed

management_group - How the waterpoint is managed

payment - What the water costs

payment_type - What the water costs

water_quality - The quality of the water

quality_group - The quality of the water

quantity - The quantity of water

quantity_group - The quantity of water

source - The source of the water

source_type - The source of the water

source_class - The source of the water

waterpoint_type - The kind of waterpoint


waterpoint_type_group - The kind of waterpoint

```
In [2]: independants = pd.read_csv("Data/Training_Set_Values.csv")
independants.head()
```

Out[2]:

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_r
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zah
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	Mal
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	Zah Nany
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825359	Sh

5 rows × 40 columns



```
In [3]: dependants = pd.read_csv("Data/Training_Set_Labels.csv")
dependants.head()
```

Out[3]:

	id	status_group
0	69572	functional
1	8776	functional
2	34310	functional
3	67743	non functional
4	19728	functional


Looks like these are our independant and dependant variables. Going to merge them together for a df.

```
In [4]: df = independants.merge(dependants, how='outer')  
df.head()
```

Out[4]:

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_r
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zah
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	Mal
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	Zah Nanyi
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825359	Sh

5 rows × 41 columns



In [5]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 59400 entries, 0 to 59399
Data columns (total 41 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    59400 non-null  int64
1   amount_tsh                           59400 non-null  float64
2   date_recorded                         59400 non-null  object
3   funder                                55765 non-null  object
4   gps_height                            59400 non-null  int64
5   installer                             55745 non-null  object
6   longitude                             59400 non-null  float64
7   latitude                             59400 non-null  float64
8   wpt_name                              59400 non-null  object
9   num_private                           59400 non-null  int64
10  basin                                 59400 non-null  object
11  subvillage                            59029 non-null  object
12  region                                59400 non-null  object
13  region_code                           59400 non-null  int64
14  district_code                         59400 non-null  int64
15  lga                                    59400 non-null  object
16  ward                                  59400 non-null  object
17  population                            59400 non-null  int64
18  public_meeting                        56066 non-null  object
19  recorded_by                           59400 non-null  object
20  scheme_management                     55523 non-null  object
21  scheme_name                           31234 non-null  object
22  permit                                56344 non-null  object
23  construction_year                     59400 non-null  int64
24  extraction_type                       59400 non-null  object
25  extraction_type_group                  59400 non-null  object
26  extraction_type_class                  59400 non-null  object
27  management                             59400 non-null  object
28  management_group                       59400 non-null  object
29  payment                                59400 non-null  object
30  payment_type                           59400 non-null  object
31  water_quality                          59400 non-null  object
32  quality_group                          59400 non-null  object
33  quantity                               59400 non-null  object
34  quantity_group                         59400 non-null  object
35  source                                 59400 non-null  object
36  source_type                            59400 non-null  object
37  source_class                           59400 non-null  object
38  waterpoint_type                        59400 non-null  object
39  waterpoint_type_group                  59400 non-null  object
40  status_group                           59400 non-null  object
dtypes: float64(3), int64(7), object(31)
memory usage: 19.0+ MB
```

Data Cleaning and Exploration

```
In [6]: ##Dropping id since it's irrelevant  
df = df.drop("id", axis=1)
```

```
In [7]: df.duplicated().sum()
```

```
Out[7]: 36
```

```
In [8]: #Dropping duplicate rows  
df.drop_duplicates(keep="first", inplace=True)
```

```
In [9]: df.duplicated().sum()
```

```
Out[9]: 0
```

```
In [10]: df.isna().sum()
```


```
Out[10]: amount_tsh          0
date_recorded              0
funder                    3635
gps_height                 0
installer                 3655
longitude                  0
latitude                   0
wpt_name                   0
num_private                0
basin                      0
subvillage                371
region                     0
region_code                0
district_code              0
lga                        0
ward                       0
population                 0
public_meeting            3314
recorded_by                0
scheme_management         3877
scheme_name               28139
permit                    3056
construction_year          0
extraction_type            0
extraction_type_group      0
extraction_type_class      0
management                 0
management_group           0
payment                    0
payment_type               0
water_quality              0
quality_group              0
quantity                   0
quantity_group             0
source                     0
source_type                0
source_class               0
waterpoint_type            0
waterpoint_type_group      0
status_group               0
dtype: int64
```

scheme_name is missing about half it's values. Dropping it

```
In [11]: df = df.drop("scheme_name", axis=1)
```

```
In [12]: plt.figure(figsize = (10,10))  
sns.scatterplot(x='longitude',y='latitude',hue='status_group',data=df)
```

```
Out[12]: <AxesSubplot:xlabel='longitude', ylabel='latitude'>
```



Looks like 0 longitude and latitude is used as a placeholder for unknown locations. Let's remove the rows


```
In [13]: df = df[df.longitude > 0]
```



```
In [14]: df.longitude.value_counts()
```

```
Out[14]: 39.090448    2
          39.086287    2
          39.086183    2
          39.098514    2
          39.093095    2
          ..
          37.579803    1
          33.196490    1
          34.017119    1
          33.788326    1
          35.005922    1
          Name: longitude, Length: 57515, dtype: int64
```

```
In [15]: plt.figure(figsize = (14,14))
sns.scatterplot(x='longitude', y='latitude', hue='status_group', data=df)
plt.title('Waterpoint Status by Geographic Location', fontsize=20)
plt.savefig('Location Map')
```



In [16]: `df.info()`


```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 57587 entries, 0 to 59399
Data columns (total 39 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   amount_tsh                            57587 non-null  float64
1   date_recorded                          57587 non-null  object
2   funder                                 53965 non-null  object
3   gps_height                             57587 non-null  int64
4   installer                              53951 non-null  object
5   longitude                              57587 non-null  float64
6   latitude                               57587 non-null  float64
7   wpt_name                               57587 non-null  object
8   num_private                            57587 non-null  int64
9   basin                                  57587 non-null  object
10  subvillage                             57216 non-null  object
11  region                                 57587 non-null  object
12  region_code                            57587 non-null  int64
13  district_code                          57587 non-null  int64
14  lga                                     57587 non-null  object
15  ward                                   57587 non-null  object
16  population                             57587 non-null  int64
17  public_meeting                         54611 non-null  object
18  recorded_by                            57587 non-null  object
19  scheme_management                      53837 non-null  object
20  permit                                 54531 non-null  object
21  construction_year                      57587 non-null  int64
22  extraction_type                         57587 non-null  object
23  extraction_type_group                  57587 non-null  object
24  extraction_type_class                  57587 non-null  object
25  management                             57587 non-null  object
26  management_group                       57587 non-null  object
27  payment                                57587 non-null  object
28  payment_type                           57587 non-null  object
29  water_quality                          57587 non-null  object
30  quality_group                          57587 non-null  object
31  quantity                               57587 non-null  object
32  quantity_group                         57587 non-null  object
33  source                                 57587 non-null  object
34  source_type                            57587 non-null  object
35  source_class                           57587 non-null  object
36  waterpoint_type                        57587 non-null  object
37  waterpoint_type_group                  57587 non-null  object
38  status_group                           57587 non-null  object
dtypes: float64(3), int64(6), object(30)
memory usage: 17.6+ MB
```

In [17]: *#date_recorded is an object. Chnaging to datetime*
`df['date_recorded'] = pd.to_datetime(df['date_recorded'])`
`df['date_recorded'] = df['date_recorded'].map(dt.datetime.toordinal)`

```
In [18]: df.date_recorded.head()
```

```
Out[18]: 0    734210
1    734933
2    734924
3    734896
4    734331
Name: date_recorded, dtype: int64
```


```
In [19]: plt.figure(figsize=(16,4))
ax = sns.countplot(x="basin", hue='status_group', data=df)
plt.xlabel('Water Basin Region', fontsize=16)
plt.title("WaterPoint Status by Basin", fontsize=20)
labels = ['non-functioning', 'functioning', 'needs repair']
plt.legend(labels)
plt.show()
```



Continuous variable cleaning

```
In [20]: df.amount_tsh.hist()
```

```
Out[20]: <AxesSubplot:>
```



```
In [21]: df.amount_tsh.value_counts(normalize=True)
```

```
Out[21]: 0.0          0.691580
          500.0        0.053866
          50.0         0.042926
          1000.0       0.025839
          20.0         0.025405
          ...
          8500.0       0.000017
          6300.0       0.000017
          220.0        0.000017
          138000.0     0.000017
          12.0         0.000017
Name: amount_tsh, Length: 98, dtype: float64
```


Most values are 0. Seems to be some outliers making the hist hard to read. Removing outliers

```
In [22]: amount_tsh_std = df.amount_tsh.mean() + df.amount_tsh.std()*3
amount_tsh_std
# Showing how many we are removing for reference
print("Outliers:", df.amount_tsh[df['amount_tsh'] > amount_tsh_std].count())
# Remove outliers from the data
df = df[df['amount_tsh'] < amount_tsh_std]
```

Outliers: 237

```
In [23]: df.amount_tsh.hist()
```

```
Out[23]: <AxesSubplot:>
```



In [24]: `df.describe()`

Out[24]:


	amount_tsh	date_recorded	gps_height	longitude	latitude	num_private	region
count	57350.000000	57350.000000	57350.000000	57350.000000	57350.000000	57350.000000	57350
mean	220.729431	734587.341883	687.53299	35.145596	-5.883289	0.484882	15
std	770.155360	335.843580	693.40072	2.608792	2.810017	12.430483	17
min	0.000000	731137.000000	-90.00000	29.607122	-11.649440	0.000000	1
25%	0.000000	734226.000000	0.00000	33.280201	-8.640322	0.000000	5
50%	0.000000	734784.000000	421.00000	35.000347	-5.168022	0.000000	12
75%	25.000000	734908.000000	1331.00000	37.231554	-3.373151	0.000000	17
max	9000.000000	735205.000000	2770.00000	40.345193	-0.998464	1776.000000	99



gps_height

In [25]: `df.gps_height.hist()`

Out[25]: <AxesSubplot:>



In [26]: `df.gps_height.describe()`

Out[26]:

count	57350.00000
mean	687.53299
std	693.40072
min	-90.00000
25%	0.00000
50%	421.00000
75%	1331.00000
max	2770.00000

Name: gps_height, dtype: float64

```
In [27]: gps_height_std = df.gps_height.mean() + df.gps_height.std()*3
gps_height_std
# Showing how many we are removing for reference
print("Outliers:", df.gps_height[df['gps_height'] > gps_height_std].count())
# Remove outliers from the data
df = df[df['gps_height'] < gps_height_std]
```

Outliers: 1

num_private


There is no description for this column.

```
In [28]: #Drop this because it is almost constant
df.num_private.value_counts(normalize=True)
```

```
Out[28]: 0      0.986922
6      0.001412
1      0.001255
5      0.000802
8      0.000802
...
180    0.000017
213    0.000017
23     0.000017
55     0.000017
94     0.000017
Name: num_private, Length: 65, dtype: float64
```

```
In [29]: df.num_private.hist()
```

Out[29]: <AxesSubplot:>



```
In [30]: df = df.drop('num_private', axis=1)
```

In []:

region_code, district_code both are prob cat

In [31]: `df.region_code.describe()`

```
Out[31]: count    57349.000000
mean       15.242341
std        17.877940
min         1.000000
25%         5.000000
50%        12.000000
75%        17.000000
max         99.000000
Name: region_code, dtype: float64
```

Population

population around the well

In [32]: `df.population.value_counts(normalize=True)`

```
Out[32]: 0      0.340651
1      0.122060
200     0.033741
150     0.032485
250     0.029259
...
363     0.000017
491     0.000017
2570    0.000017
587     0.000017
1439    0.000017
Name: population, Length: 1042, dtype: float64
```

A lot of 0 values. Abandoned? Surely not.

In [33]: `df.population.describe()`

```
Out[33]: count    57349.000000
mean       185.269630
std        478.157306
min         0.000000
25%         0.000000
50%        35.000000
75%       230.000000
max       30500.000000
Name: population, dtype: float64
```



```
In [34]: df.population.median()
```

```
Out[34]: 35.0
```

I think I'll replace the 0 values with the median. Will do this after split

```
In [35]: df.population.replace(0,df.population.median(axis=0),inplace=True)
```

```
In [36]: df.population.describe()
```

```
Out[36]: count      57349.000000
mean         197.192418
std          473.805467
min           1.000000
25%          35.000000
50%          35.000000
75%         230.000000
max        30500.000000
Name: population, dtype: float64
```

```
In [35]: population_std = df.population.mean() + df.population.std()*3
population_std
# Showing how many we are removing for reference
print("Outliers:", df.population[df['population'] > population_std].count())
# Remove outliers from the data
df = df[df['population'] < population_std]
```

```
Outliers: 702
```

construction_year

```
In [36]: df.construction_year.value_counts(normalize=True)
```

```
Out[36]: 0      0.332921
2010     0.045104
2008     0.045033
2009     0.043639
2000     0.036154
2007     0.027574
2006     0.025156
2003     0.022243
2011     0.021572
2004     0.019471
2012     0.018783
1978     0.018130
2002     0.017883
2005     0.017494
1995     0.016929
1999     0.016929
1998     0.016629
1990     0.016629
1985     0.016312
1980     0.014158
1996     0.013787
1984     0.013416
1982     0.012869
1994     0.012587
1972     0.012357
1974     0.011580
1997     0.011210
1992     0.011033
1993     0.010451
2001     0.009356
1988     0.009091
1983     0.008438
1975     0.007573
1986     0.007538
1976     0.007220
1970     0.007167
1991     0.005649
1989     0.005490
1987     0.005190
1981     0.004148
1977     0.003548
1979     0.003354
1973     0.003195
2013     0.003107
1971     0.002507
1960     0.001801
1967     0.001553
1963     0.001465
1968     0.001324
1969     0.001024
1964     0.000706
1962     0.000530
1961     0.000371
1965     0.000335
```

```
1966      0.000282  
Name: construction_year, dtype: float64
```

Once again a lot of 0's here


```
In [37]: year = df[df['construction_year'] != 0]
```

```
In [38]: year.construction_year.describe()
```

```
Out[38]: count      37788.000000  
mean        1996.768233  
std          12.500440  
min          1960.000000  
25%          1987.000000  
50%          2000.000000  
75%          2008.000000  
max          2013.000000  
Name: construction_year, dtype: float64
```

```
In [39]: year.construction_year.hist(bins=20)
```

```
Out[39]: <AxesSubplot:>
```



Categorical Cleaning

```
In [40]: df.isna().sum()
```

```
Out[40]: amount_tsh          0
         date_recorded       0
         funder             3588
         gps_height          0
         installer          3602
         longitude          0
         latitude           0
         wpt_name           0
         basin              0
         subvillage         371
         region             0
         region_code        0
         district_code      0
         lga                0
         ward               0
         population         0
         public_meeting     2929
         recorded_by        0
         scheme_management  3679
         permit             2997
         construction_year  0
         extraction_type    0
         extraction_type_group 0
         extraction_type_class 0
         management         0
         management_group   0
         payment            0
         payment_type       0
         water_quality      0
         quality_group      0
         quantity           0
         quantity_group     0
         source             0
         source_type        0
         source_class       0
         waterpoint_type    0
         waterpoint_type_group 0
         status_group       0
         dtype: int64
```

funder

```
In [41]: df.funder.describe()
```

```
Out[41]: count          53059
         unique          1823
         top      Government Of Tanzania
         freq              8784
         Name: funder, dtype: object
```

```
In [42]: #tons of unique/missing values. Dropping and focusing on others
df =df.drop('funder', axis=1)
```

```
In [43]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 56647 entries, 0 to 59399
Data columns (total 37 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   amount_tsh                           56647 non-null  float64
1   date_recorded                        56647 non-null  int64
2   gps_height                           56647 non-null  int64
3   installer                            53045 non-null  object
4   longitude                            56647 non-null  float64
5   latitude                             56647 non-null  float64
6   wpt_name                             56647 non-null  object
7   basin                               56647 non-null  object
8   subvillage                           56276 non-null  object
9   region                              56647 non-null  object
10  region_code                          56647 non-null  int64
11  district_code                       56647 non-null  int64
12  lga                                  56647 non-null  object
13  ward                                56647 non-null  object
14  population                           56647 non-null  int64
15  public_meeting                      53718 non-null  object
16  recorded_by                         56647 non-null  object
17  scheme_management                   52968 non-null  object
18  permit                             53650 non-null  object
19  construction_year                   56647 non-null  int64
20  extraction_type                     56647 non-null  object
21  extraction_type_group                56647 non-null  object
22  extraction_type_class                56647 non-null  object
23  management                           56647 non-null  object
24  management_group                    56647 non-null  object
25  payment                             56647 non-null  object
26  payment_type                        56647 non-null  object
27  water_quality                       56647 non-null  object
28  quality_group                       56647 non-null  object
29  quantity                            56647 non-null  object
30  quantity_group                      56647 non-null  object
31  source                              56647 non-null  object
32  source_type                         56647 non-null  object
33  source_class                        56647 non-null  object
34  waterpoint_type                     56647 non-null  object
35  waterpoint_type_group               56647 non-null  object
36  status_group                        56647 non-null  object
dtypes: float64(3), int64(6), object(28)
memory usage: 16.4+ MB
```

Installer

```
In [44]: df.installer.describe()
```

```
Out[44]: count      53045  
         unique      2069  
         top         DWE  
         freq      16051  
         Name: installer, dtype: object
```

```
In [45]: #Dropping for same reason.  
df = df.drop('installer', axis=1)
```

wpt_name

This is just the name of the waterpoint. Dropping because should be irrelevant

```
In [46]: df = df.drop('wpt_name', axis=1)
```

Basin


```
In [47]: df.basin.describe()
```

```
Out[47]: count      56647  
         unique        9  
         top      Pangani  
         freq      8805  
         Name: basin, dtype: object
```

```
In [48]: df.basin.value_counts(normalize=True)
```


```
Out[48]: Pangani                0.155436  
         Lake Victoria          0.148393  
         Rufiji                 0.139054  
         Internal               0.135541  
         Lake Tanganyika        0.109732  
         Wami / Ruvu            0.102742  
         Lake Nyasa             0.088760  
         Ruvuma / Southern Coast 0.077798  
         Lake Rukwa             0.042544  
         Name: basin, dtype: float64
```

```
In [49]: df.basin.value_counts().plot(kind='barh')
plt.xlabel('Number of Waterpoints', fontsize=16)
plt.ylabel('Region', fontsize=16)
plt.title("WaterPoints by Basin", fontsize=20)
plt.show()
```



In [50]:

```
plt.figure(figsize=(16,4))
ax = sns.countplot(x="basin", hue='status_group', data=df)
plt.xlabel('Water Basin Region', fontsize=16)
plt.ylabel('Number of WaterPoints', fontsize=16)
plt.title("WaterPoint Status by Basin", fontsize=20)
labels = ['non-functioning', 'functioning', 'needs repair']
plt.legend(labels)
plt.show()
```



subvillage

In [51]: `df.subvillage.describe()`

```
Out[51]: count      56276
unique     18348
top        Shuleni
freq         482
Name: subvillage, dtype: object
```

In [52]: `df = df.drop('subvillage', axis=1)`

Region

In [53]: `df.region.describe()`

```
Out[53]: count      56647
unique        21
top         Iringa
freq        5249
Name: region, dtype: object
```

```
In [54]: #already have region code which is the same thing. Dropping
df = df.drop('region', axis=1)
```

Looking at the data a bit, I'm going to drop lga and ward as well for similar reasons. Too many geographic features


```
In [55]: df = df.drop(columns='lga', axis=1)
df = df.drop(columns='ward', axis=1)
```

public_meeting

Whether the waterpoint is open to the public

some missing values

```
In [56]: df.public_meeting.describe()
```

```
Out[56]: count      53718
unique         2
top            True
freq         48953
Name: public_meeting, dtype: object
```

```
In [57]: df.public_meeting.value_counts(normalize=True)
```

```
Out[57]: True      0.911296
False    0.088704
Name: public_meeting, dtype: float64
```

Almost all are open to public. Going to fill missing values with true

```
In [58]: df['public_meeting'].fillna(True, inplace=True)
```

recorded_by

```
In [59]: df.recorded_by.describe()
```

```
Out[59]: count      56647
unique         1
top      GeoData Consultants Ltd
freq      56647
Name: recorded_by, dtype: object
```

All recorded by the same group. Removing since this won't be useful

```
In [60]: df = df.drop('recorded_by', axis=1)
```

Scheme_management

```
In [61]: df.scheme_management.describe()
```

```
Out[61]: count      52968
         unique        12
         top         VWC
         freq      35613
         Name: scheme_management, dtype: object
```

```
In [62]: df.scheme_management.value_counts()
```

```
Out[62]: VWC      35613
         WUG      4205
         Water authority  3081
         WUA      2847
         Water Board  2717
         Parastatal  1579
         Private operator  1034
         Company    1034
         Other      691
         SWC        96
         Trust      70
         None        1
         Name: scheme_management, dtype: int64
```


```
In [63]: #Dropping for same reason as many above
         df = df.drop('scheme_management', axis=1)
```

permit

```
In [64]: df.permit.value_counts()
```

```
Out[64]: True      37526
         False    16124
         Name: permit, dtype: int64
```

```
In [65]: plt.figure(figsize=(16,4))
         ax = sns.countplot(x="permit", hue='status_group', data=df)
         labels = ['non-functioning', 'functioning', 'needs repair']
         plt.legend(labels)
         plt.show()
```



Weird. I expected those with permits to be more regulated than those without

There are some missing values here. Filling them with false.

```
In [66]: df['permit'].fillna(False, inplace=True)
```

extraction

These 3 seem to have similar data.

```
In [67]: df.extraction_type.value_counts()
```

```
Out[67]: gravity                26389
nira/tanira                    7269
other                          6050
submersible                    4535
swn 80                         3404
mono                          2767
india mark ii                  2206
afridev                        1609
ksb                            1331
other - rope pump              444
other - swn 81                 224
windmill                      107
cemo                           90
india mark iii                 87
other - play pump              81
climax                        32
walimi                         20
other - mkulima/shinyanga      2
Name: extraction_type, dtype: int64
```

```
In [68]: df.extraction_type_group.value_counts()
```

```
Out[68]: gravity                26389
nira/tanira                    7269
other                          6050
submersible                    5866
swn 80                         3404
mono                          2767
india mark ii                  2206
afridev                        1609
rope pump                      444
other handpump                 327
other motorpump               122
wind-powered                   107
india mark iii                 87
Name: extraction_type_group, dtype: int64
```

```
In [69]: df.extraction_type_class.value_counts()
```

```
Out[69]: gravity          26389
handpump          14902
other             6050
submersible       5866
motorpump         2889
rope pump          444
wind-powered      107
Name: extraction_type_class, dtype: int64
```

Going to keep class and drop the other two.


```
In [70]: df = df.drop(columns=['extraction_type', 'extraction_type_group'], axis=1)
```

```
In [71]: #putting low cardinality into other group
df.extraction_type_class = df.extraction_type_class.replace(to_replace =
                                                             ['rope pump',
                                                             'wind-powered'],
                                                             value = 'other')

df.extraction_type_class.value_counts()
```

```
Out[71]: gravity          26389
handpump          14902
other             6601
submersible       5866
motorpump         2889
Name: extraction_type_class, dtype: int64
```

```
In [72]: plt.figure(figsize=(10,4))
ax = sns.countplot(x="extraction_type_class", hue='status_group', data=df)
plt.xlabel('Extraction Type', fontsize=16)
plt.ylabel('Number of WaterPoints', fontsize=16)
plt.title("Status of Different Types of Waterpoints", fontsize=20)
labels = ['non-functioning', 'functioning', 'needs repair']
plt.legend(labels)
plt.show()
```



management

```
In [73]: df.management.value_counts()
```

```
Out[73]: vwc          39119  
wug          5498  
water board  2888  
wua          2509  
private operator  1923  
parastatal   1660  
water authority  881  
other        791  
company      663  
unknown      541  
other - school  98  
trust        76  
Name: management, dtype: int64
```

```
In [74]: df.management_group.value_counts()
```

```
Out[74]: user-group    50014  
commercial    3543  
parastatal    1660  
other         889  
unknown       541  
Name: management_group, dtype: int64
```

```
In [75]: df = df.drop('management', axis=1)
```

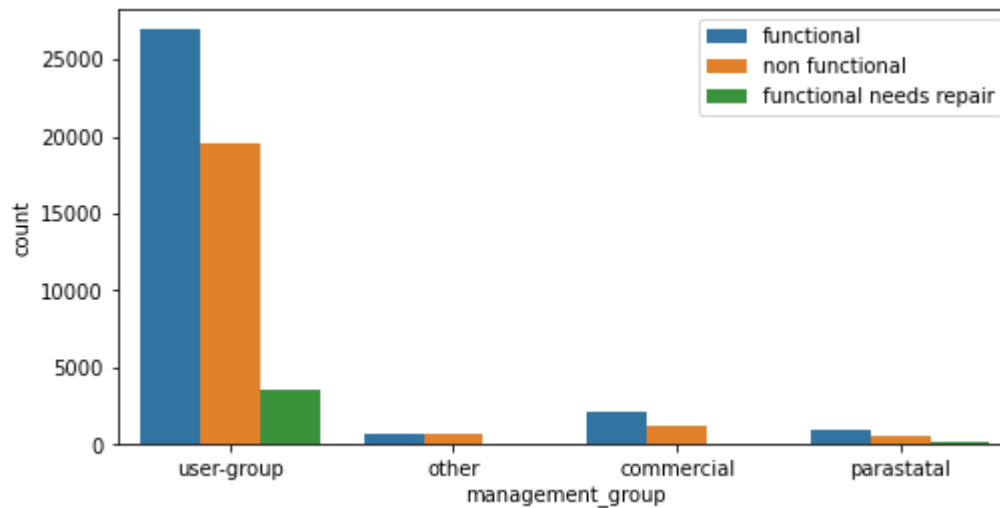
```
In [76]: df.management_group = df.management_group.replace(to_replace = ['unknown'],  
                                                             value = 'other')  
df.management_group.value_counts()
```

```
Out[76]: user-group    50014  
commercial    3543  
parastatal    1660  
other        1430  
Name: management_group, dtype: int64
```

In [77]:

```
plt.figure(figsize=(8,4))  
ax = sns.countplot(x="management_group", hue='status_group', data=df)  
plt.legend()
```

Out[77]: <matplotlib.legend.Legend at 0x237bf68b760>



Payment

In [78]: df.payment.value_counts()

Out[78]:

never pay	24097
pay per bucket	8707
pay monthly	8059
unknown	7582
pay when scheme fails	3804
pay annually	3506
other	892

Name: payment, dtype: int64

```
In [79]: df.payment_type.value_counts()
```

```
Out[79]: never pay      24097
per bucket    8707
monthly      8059
unknown      7582
on failure    3804
annually     3506
other        892
Name: payment_type, dtype: int64
```

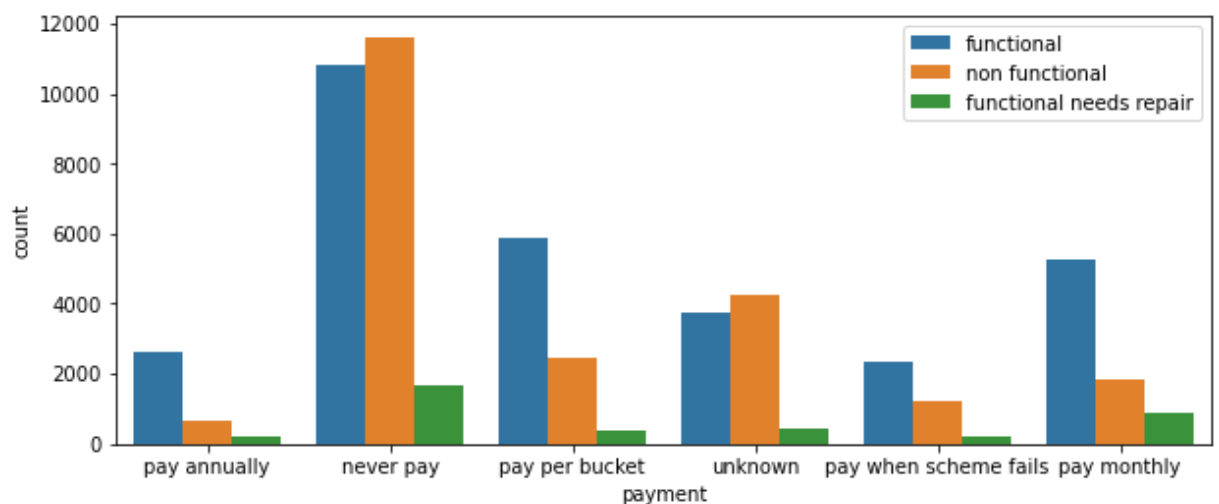
```
In [80]: #Looks liek a duplicate column. Removing extra
df = df.drop('payment_type', axis=1)
```

```
In [81]: # group 'unknowns' and 'other' together to reduce bins
df.payment = df.payment.replace(to_replace = ['other'],
                                value = 'unknown')
df.payment.value_counts()
```

```
Out[81]: never pay      24097
pay per bucket    8707
unknown          8474
pay monthly      8059
pay when scheme fails  3804
pay annually     3506
Name: payment, dtype: int64
```

```
In [82]: plt.figure(figsize=(10,4))
ax = sns.countplot(x="payment", hue='status_group', data=df)
plt.legend()
```

```
Out[82]: <matplotlib.legend.Legend at 0x237bf404820>
```



Water Quality

```
In [83]: df.quality_group.value_counts()
```

```
Out[83]: good          48664
salty           4862
unknown        1646
milky           798
colored         473
fluoride        204
Name: quality_group, dtype: int64
```

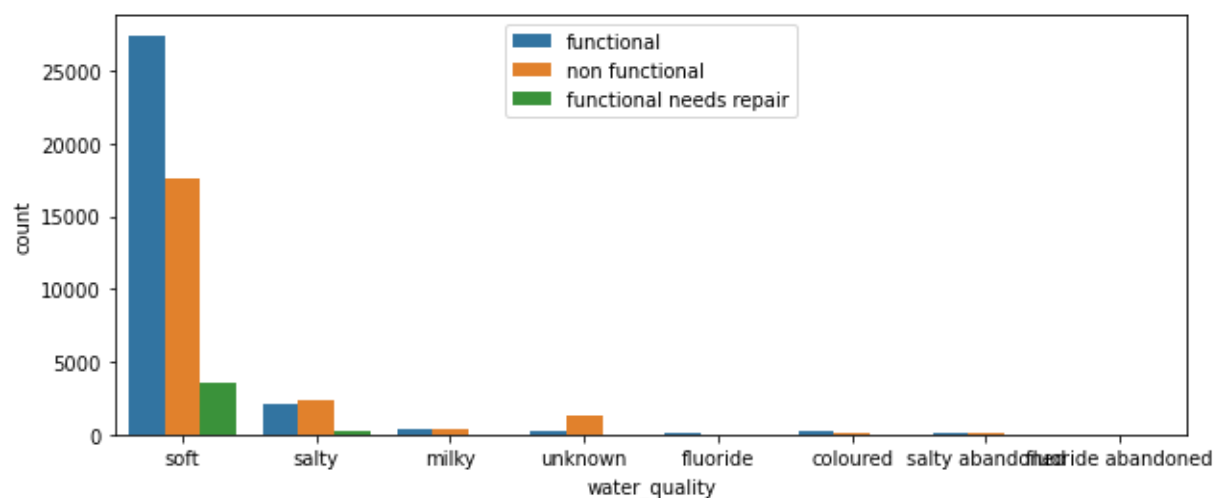
```
In [ ]:
```

These are almost identical. Just notes the abandoned. Dropping former

```
In [84]: df = df.drop('quality_group', axis=1)
```

```
In [85]: plt.figure(figsize=(10,4))
ax = sns.countplot(x="water_quality", hue='status_group', data=df)
plt.legend()
```

```
Out[85]: <matplotlib.legend.Legend at 0x237bf785e80>
```



```
In [86]: # Group all to make it wither good or bad water
df.water_quality = df.water_quality.replace(to_replace = ['salty', 'milky', 'unkr
                                                    'fluoride', 'coloured', 'salty abandon
                                                    'fluoride abandoned'],
                                                    value = 'other')

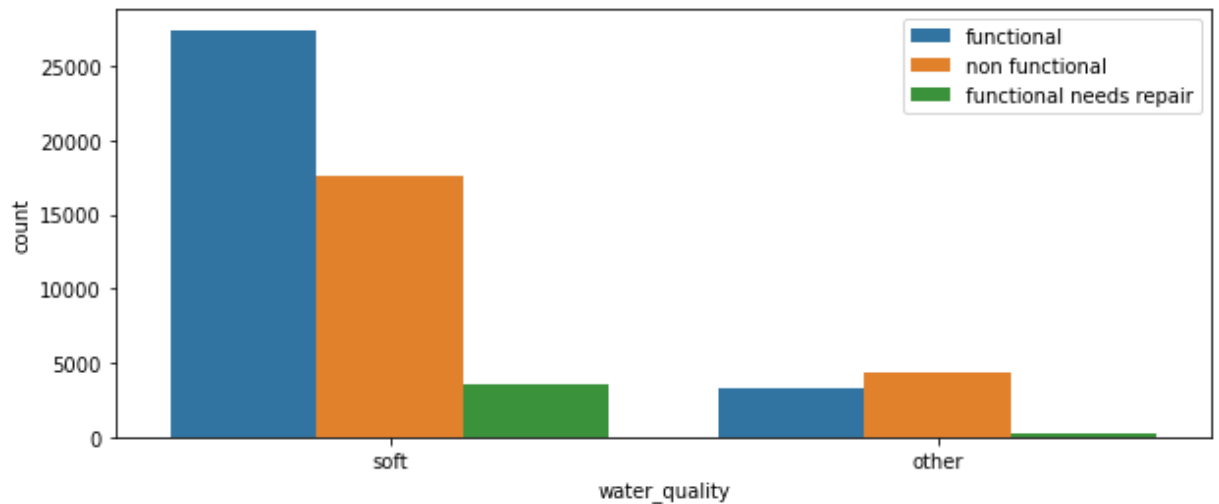
df.water_quality.value_counts()
```

```
Out[86]: soft          48664
other           7983
Name: water_quality, dtype: int64
```



```
In [87]: plt.figure(figsize=(10,4))
ax = sns.countplot(x="water_quality", hue='status_group', data=df)
plt.legend()
```

Out[87]: <matplotlib.legend.Legend at 0x237c0667dc0>



Quantity

```
In [88]: df.quantity.value_counts()
```

Out[88]:

enough	31684
insufficient	14302
dry	5934
seasonal	3965
unknown	762

Name: quantity, dtype: int64

```
In [89]: df.quantity_group.value_counts()
```

Out[89]:

enough	31684
insufficient	14302
dry	5934
seasonal	3965
unknown	762

Name: quantity_group, dtype: int64

```
In [90]: #duplicated columns
df = df.drop('quantity_group', axis=1)
```


```
In [91]: # changing groups to either sufficient or not
df.quantity = df.quantity.replace(to_replace = ['unknown', 'dry',
                                                'seasonal'],
                                  value = 'insufficient')

df.quantity.value_counts()
```

```
Out[91]: enough          31684
insufficient    24963
Name: quantity, dtype: int64
```

```
In [92]: plt.figure(figsize=(10,4))
ax = sns.countplot(x="quantity", hue='status_group', data=df)
plt.legend()
```

```
Out[92]: <matplotlib.legend.Legend at 0x237c06c0d30>
```



Source

```
In [93]: df.source.value_counts()
```

```
Out[93]: spring          16833
shallow well    15288
machine dbh     10555
river           9409
rainwater harvesting  2194
hand dtw        864
dam             621
lake           620
other          202
unknown        61
Name: source, dtype: int64
```

```
In [94]: df.source_class.value_counts()
```

```
Out[94]: groundwater    43540
surface      12844
unknown      263
Name: source_class, dtype: int64
```

```
In [95]: df.source_type.value_counts()
```

```
Out[95]: spring                16833
         shallow well          15288
         borehole              11419
         river/lake            10029
         rainwater harvesting   2194
         dam                   621
         other                  263
         Name: source_type, dtype: int64
```

Looks like these three are redundant info dropping two

```
In [96]: df = df.drop(columns=['source', 'source_type'], axis=1)
```

Most of the sources are from groundwater. Very little data in unknown cat

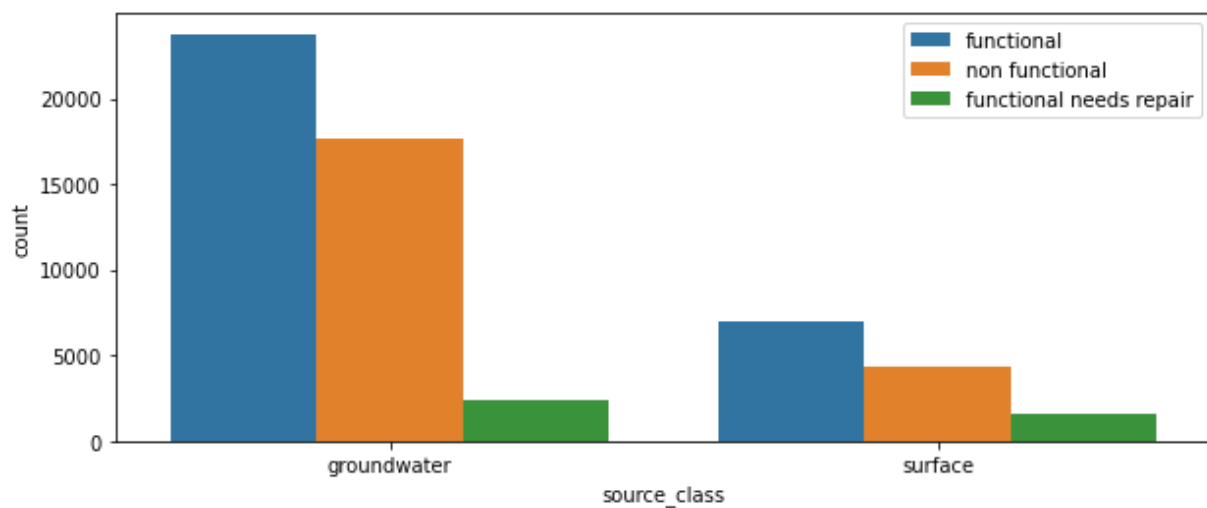
Putting those values under groundwater to reduce bins

```
In [97]: df.source_class = df.source_class.replace(to_replace = 'unknown',
                                                    value = 'groundwater')
         df.source_class.value_counts()
```

```
Out[97]: groundwater    43803
         surface        12844
         Name: source_class, dtype: int64
```

```
In [98]: plt.figure(figsize=(10,4))
         ax = sns.countplot(x="source_class", hue='status_group', data=df)
         plt.legend()
```

```
Out[98]: <matplotlib.legend.Legend at 0x237c0855ca0>
```



Waterpoint Type

```
In [99]: df.waterpoint_type_group.value_counts()
```

```
Out[99]: communal standpipe    33811
hand pump                    15882
other                       6076
improved spring              767
cattle trough                105
dam                          6
Name: waterpoint_type_group, dtype: int64
```


```
In [100]: df.waterpoint_type.value_counts()
```

```
Out[100]: communal standpipe    28012
hand pump                    15882
other                       6076
communal standpipe multiple    5799
improved spring              767
cattle trough                105
dam                          6
Name: waterpoint_type, dtype: int64
```

```
In [101]: #dropping similar columns
df = df.drop('waterpoint_type_group', axis=1)
```

```
In [102]: plt.figure(figsize=(14,4))
ax = sns.countplot(x="waterpoint_type", hue='status_group', data=df)
plt.legend()
```

```
Out[102]: <matplotlib.legend.Legend at 0x237c1450070>
```



I'm going to group some of these together

```
In [103]: # all standpipes will be in a group
df.waterpoint_type = df.waterpoint_type.replace(to_replace = 'communal standpipe',
                                                  value = 'communal standpipe')

df.waterpoint_type.value_counts()
```

```
Out[103]: communal standpipe    33811
hand pump                    15882
other                       6076
improved spring              767
cattle trough                105
dam                          6
Name: waterpoint_type, dtype: int64
```


```
In [104]: # Putting all the lower cardinality into the other group leaving only 3 columns
df.waterpoint_type = df.waterpoint_type.replace(to_replace = ['improved spring',
                                                             'cattle trough', 'dam'],
                                                  value = 'other')

df.waterpoint_type.value_counts()
```

```
Out[104]: communal standpipe    33811
hand pump                    15882
other                       6954
Name: waterpoint_type, dtype: int64
```

```
In [105]: plt.figure(figsize=(14,4))
ax = sns.countplot(x="waterpoint_type", hue='status_group', data=df)
plt.legend()
```

```
Out[105]: <matplotlib.legend.Legend at 0x237c14502b0>
```



In [106]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 56647 entries, 0 to 59399
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   amount_tsh            56647 non-null  float64
1   date_recorded          56647 non-null  int64
2   gps_height             56647 non-null  int64
3   longitude              56647 non-null  float64
4   latitude               56647 non-null  float64
5   basin                  56647 non-null  object
6   region_code            56647 non-null  int64
7   district_code          56647 non-null  int64
8   population             56647 non-null  int64
9   public_meeting         56647 non-null  bool
10  permit                 56647 non-null  bool
11  construction_year      56647 non-null  int64
12  extraction_type_class   56647 non-null  object
13  management_group        56647 non-null  object
14  payment                 56647 non-null  object
15  water_quality           56647 non-null  object
16  quantity                56647 non-null  object
17  source_class            56647 non-null  object
18  waterpoint_type         56647 non-null  object
19  status_group            56647 non-null  object
dtypes: bool(2), float64(3), int64(6), object(9)
memory usage: 10.8+ MB
```

Cleaned Data

One Hot Encoding Categoricals

```
In [107]: cat_dummies = ['basin', 'public_meeting', 'permit',
                        'extraction_type_class', 'management_group', 'payment', 'water_quality',
                        'quantity', 'source_class', 'waterpoint_type']
```

```
In [108]: enc = OneHotEncoder(handle_unknown='ignore')
# passing bridge-types-cat column (label encoded values of bridge_types)
enc_df = pd.DataFrame(enc.fit_transform(df[['basin', 'public_meeting', 'permit',
      'extraction_type_class', 'management_group', 'payment', 'water_quality',
      'quantity', 'source_class', 'waterpoint_type']]).toarray())
# merge with main df bridge_df on key values
#bridge_df = bridge_df.join(enc_df)
#bridge_df
```


```
In [109]: enc.get_feature_names()
enc_df.columns = enc.get_feature_names()
```

```
In [110]: enc_df.head()
```

Out[110]:

	x0_Internal	x0_Lake Nyasa	x0_Lake Rukwa	x0_Lake Tanganyika	x0_Lake Victoria	x0_Pangani	x0_Rufiji	x0_Ruvuma / Southern Coast	x0_Wa / Ru
0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	(
1	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	(
2	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	(
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	(
4	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	(

5 rows × 37 columns



```
In [111]: enc_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 56647 entries, 0 to 56646
Data columns (total 37 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   x0_Internal                          56647 non-null  float64
1   x0_Lake Nyasa                        56647 non-null  float64
2   x0_Lake Rukwa                       56647 non-null  float64
3   x0_Lake Tanganyika                  56647 non-null  float64
4   x0_Lake Victoria                    56647 non-null  float64
5   x0_Pangani                          56647 non-null  float64
6   x0_Rufiji                           56647 non-null  float64
7   x0_Ruvuma / Southern Coast          56647 non-null  float64
8   x0_Wami / Ruvu                      56647 non-null  float64
9   x1_False                            56647 non-null  float64
10  x1_True                             56647 non-null  float64
11  x2_False                            56647 non-null  float64
12  x2_True                             56647 non-null  float64
13  x3_gravity                          56647 non-null  float64
14  x3_handpump                         56647 non-null  float64
15  x3_motorpump                       56647 non-null  float64
16  x3_other                            56647 non-null  float64
17  x3_submersible                      56647 non-null  float64
18  x4_commercial                       56647 non-null  float64
19  x4_other                            56647 non-null  float64
20  x4_parastatal                       56647 non-null  float64
21  x4_user-group                       56647 non-null  float64
22  x5_never pay                        56647 non-null  float64
23  x5_pay annually                     56647 non-null  float64
24  x5_pay monthly                      56647 non-null  float64
25  x5_pay per bucket                   56647 non-null  float64
26  x5_pay when scheme fails            56647 non-null  float64
27  x5_unknown                          56647 non-null  float64
28  x6_other                            56647 non-null  float64
29  x6_soft                             56647 non-null  float64
30  x7_enough                           56647 non-null  float64
31  x7_insufficient                     56647 non-null  float64
32  x8_groundwater                     56647 non-null  float64
33  x8_surface                          56647 non-null  float64
34  x9_communal standpipe               56647 non-null  float64
35  x9_hand pump                        56647 non-null  float64
36  x9_other                            56647 non-null  float64
dtypes: float64(37)
memory usage: 16.0 MB
```

```
In [112]: df2 = df.drop(columns=['basin', 'public_meeting', 'permit',
                                'extraction_type_class', 'management_group', 'payment', 'water_quality',
                                'quantity', 'source_class', 'waterpoint_type'], axis=1)
```


In [113]: df2.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 56647 entries, 0 to 59399
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   amount_tsh            56647 non-null  float64
1   date_recorded          56647 non-null  int64
2   gps_height             56647 non-null  int64
3   longitude              56647 non-null  float64
4   latitude               56647 non-null  float64
5   region_code            56647 non-null  int64
6   district_code          56647 non-null  int64
7   population             56647 non-null  int64
8   construction_year      56647 non-null  int64
9   status_group           56647 non-null  object
dtypes: float64(3), int64(6), object(1)
memory usage: 7.3+ MB
```

In [114]: df3 = df2.join(enc_df, how='left')

In [115]: df3.columns

```
Out[115]: Index(['amount_tsh', 'date_recorded', 'gps_height', 'longitude', 'latitude',
               'region_code', 'district_code', 'population', 'construction_year',
               'status_group', 'x0_Internal', 'x0_Lake Nyasa', 'x0_Lake Rukwa',
               'x0_Lake Tanganyika', 'x0_Lake Victoria', 'x0_Pangani', 'x0_Rufiji',
               'x0_Ruvuma / Southern Coast', 'x0_Wami / Ruvu', 'x1_False', 'x1_True',
               'x2_False', 'x2_True', 'x3_gravity', 'x3_handpump', 'x3_motorpump',
               'x3_other', 'x3_submersible', 'x4_commercial', 'x4_other',
               'x4_parastatal', 'x4_user-group', 'x5_never pay', 'x5_pay annually',
               'x5_pay monthly', 'x5_pay per bucket', 'x5_pay when scheme fails',
               'x5_unknown', 'x6_other', 'x6_soft', 'x7_enough', 'x7_insufficient',
               'x8_groundwater', 'x8_surface', 'x9_communal standpipe', 'x9_hand pump',
               'x9_other'],
              dtype='object')
```

In [116]: df3.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 56647 entries, 0 to 59399
Data columns (total 47 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   amount_tsh                           56647 non-null  float64
1   date_recorded                        56647 non-null  int64
2   gps_height                           56647 non-null  int64
3   longitude                            56647 non-null  float64
4   latitude                             56647 non-null  float64
5   region_code                          56647 non-null  int64
6   district_code                       56647 non-null  int64
7   population                           56647 non-null  int64
8   construction_year                   56647 non-null  int64
9   status_group                         56647 non-null  object
10  x0_Internal                          54030 non-null  float64
11  x0_Lake Nyasa                        54030 non-null  float64
12  x0_Lake Rukwa                       54030 non-null  float64
13  x0_Lake Tanganyika                  54030 non-null  float64
14  x0_Lake Victoria                    54030 non-null  float64
15  x0_Pangani                          54030 non-null  float64
16  x0_Rufiji                           54030 non-null  float64
17  x0_Ruvuma / Southern Coast          54030 non-null  float64
18  x0_Wami / Ruvu                      54030 non-null  float64
19  x1_False                            54030 non-null  float64
20  x1_True                             54030 non-null  float64
21  x2_False                            54030 non-null  float64
22  x2_True                             54030 non-null  float64
23  x3_gravity                          54030 non-null  float64
24  x3_handpump                         54030 non-null  float64
25  x3_motorpump                       54030 non-null  float64
26  x3_other                            54030 non-null  float64
27  x3_submersible                      54030 non-null  float64
28  x4_commercial                       54030 non-null  float64
29  x4_other                            54030 non-null  float64
30  x4_parastatal                       54030 non-null  float64
31  x4_user-group                       54030 non-null  float64
32  x5_never pay                        54030 non-null  float64
33  x5_pay annually                     54030 non-null  float64
34  x5_pay monthly                      54030 non-null  float64
35  x5_pay per bucket                   54030 non-null  float64
36  x5_pay when scheme fails            54030 non-null  float64
37  x5_unknown                          54030 non-null  float64
38  x6_other                            54030 non-null  float64
39  x6_soft                             54030 non-null  float64
40  x7_enough                           54030 non-null  float64
41  x7_insufficient                     54030 non-null  float64
42  x8_groundwater                      54030 non-null  float64
43  x8_surface                          54030 non-null  float64
44  x9_communal standpipe               54030 non-null  float64
45  x9_hand pump                        54030 non-null  float64
46  x9_other                            54030 non-null  float64
dtypes: float64(40), int64(6), object(1)
memory usage: 23.2+ MB
```

```
In [117]: df3 = df3.fillna(0.0)
```

```
In [155]: df4 = df3[df3.status_group != 'functional needs repair']  
#decided to drop the needs repair status. There are far less of these, and these
```

```
In [119]: df4.status_group.value_counts()
```

```
Out[119]: functional      30752  
non functional    22002  
Name: status_group, dtype: int64
```

Split data into test and train

```
In [120]: y = df4.status_group  
X = df4.drop('status_group', axis=1)
```

```
In [121]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_  
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[121]: ((39565, 46), (13189, 46), (39565,), (13189,))
```

Models

Random Forest

After tinkering with this model quite a bit, these are the hyperparameter we ended up using.

```
In [202]: forest_clf = RandomForestClassifier(random_state=42, max_depth=15, n_estimators=500,  
min_samples_leaf=2,  
min_samples_split=2)  
forest_model = forest_clf.fit(X_train, y_train)  
  
forest_training_preds = forest_clf.predict(X_train)  
forest_training_accuracy = accuracy_score(y_train, forest_training_preds)  
  
forest_val_preds = forest_clf.predict(X_test)  
forest_val_accuracy = accuracy_score(y_test, forest_val_preds)
```

Forest Training Accuracy:
Forest Validation accuracy:

```
In [207]: print(classification_report(y_test, forest_val_preds))
          #test report
```

	precision	recall	f1-score	support
functional	0.72	0.88	0.79	7754
non functional	0.75	0.52	0.61	5603
accuracy			0.73	13357
macro avg	0.73	0.70	0.70	13357
weighted avg	0.73	0.73	0.72	13357

```
In [208]: print(classification_report(y_train, forest_training_preds))
          #training report
```


	precision	recall	f1-score	support
functional	0.81	0.95	0.87	23428
non functional	0.90	0.69	0.78	16642
accuracy			0.84	40070
macro avg	0.86	0.82	0.83	40070
weighted avg	0.85	0.84	0.84	40070

It does look like our model is overfitting a bit, but not too much. Accuracy is 73% vs the 85% from the training

```
In [205]: print(confusion_matrix(y_test, forest_val_preds))
```

```
[[6797  957]
 [2693 2910]]
```

```
In [206]: sorted_idx = forest_clf.feature_importances_.argsort()
plt.figure(figsize=(10,12))
plt.barh(X.columns[sorted_idx], forest_clf.feature_importances_[sorted_idx])
plt.ylabel('Feature', fontsize=16)
plt.title("Feature Importance", fontsize=20)
plt.xlabel("Random Forest Feature Importance")
plt.savefig('Feature Importance')
```



In []:

In [179]: *# Forst Model with GridSearch parameters*

```
param_grid_2 = {
    'max_depth': [5, 10, 30, None],
    'min_samples_split': [2, 3],
    'min_samples_leaf': [1, 2, 5],
    'n_estimators': [10, 25, 100],
}
```

In [198]: *# GridSearch Classifier**# This time I'm going to try criterion='entropy'*

forest_clf = RandomForestClassifier(random_state=42)

grid_clf = GridSearchCV(forest_clf, param_grid_2, scoring='accuracy', cv=3, n_jobs=4)

grid_clf.fit(X_train, y_train)

best_parameters = grid_clf.best_params_

print("Grid Search found the following optimal parameters: ")

for param_name in sorted(best_parameters.keys()):

print("%s: %r" % (param_name, best_parameters[param_name]))

training_preds_forest = grid_clf.predict(X_train)

training_accuracy_forest = accuracy_score(y_train, training_preds_forest)

val_preds_forest = grid_clf.predict(X_test)

val_accuracy_forest = accuracy_score(y_test, val_preds_forest)

print("")

print("Training Accuracy: {:.4}%".format(training_accuracy_forest * 100))

print("Validation accuracy: {:.4}%".format(val_accuracy_forest * 100))

Grid Search found the following optimal parameters:

max_depth: None

min_samples_leaf: 2

min_samples_split: 2

n_estimators: 100

Training Accuracy: 96.94%

Validation accuracy: 72.97%

```
In [182]: # Classification report
print(classification_report(y_test, val_preds_forest))
```

	precision	recall	f1-score	support
functional	0.73	0.84	0.78	7754
non functional	0.73	0.57	0.64	5603
accuracy			0.73	13357
macro avg	0.73	0.71	0.71	13357
weighted avg	0.73	0.73	0.72	13357

```
In [183]: print(classification_report(y_train, training_preds_forest))
```

	precision	recall	f1-score	support
functional	0.96	0.99	0.97	23428
non functional	0.98	0.94	0.96	16642
accuracy			0.97	40070
macro avg	0.97	0.97	0.97	40070
weighted avg	0.97	0.97	0.97	40070

After running several different grid searches, I still was having trouble with overfitting. I decided to go with the previous hyperparameters instead.

XGBoost

```
In [122]: # XGB classifier
xgb_clf = xgb.XGBClassifier()
xgb_clf.fit(X_train, y_train)

xgb_training_preds = xgb_clf.predict(X_train)
xgb_training_accuracy = accuracy_score(y_train, xgb_training_preds)

xgb_val_preds = xgb_clf.predict(X_test)
xgb_val_accuracy = accuracy_score(y_test, xgb_val_preds)

print("XGB Training Accuracy: {:.4}%".format(xgb_training_accuracy * 100))
print("XGB Validation accuracy: {:.4}%".format(xgb_val_accuracy * 100))
```

XGB Training Accuracy: 81.52%
XGB Validation accuracy: 74.65%

```
In [123]: print(classification_report(y_test, xgb_val_preds))
```


	precision	recall	f1-score	support
functional	0.76	0.83	0.79	7694
non functional	0.73	0.63	0.67	5495
accuracy			0.75	13189
macro avg	0.74	0.73	0.73	13189
weighted avg	0.74	0.75	0.74	13189

```
In [124]: print(classification_report(y_train, xgb_training_preds))
```

	precision	recall	f1-score	support
functional	0.81	0.89	0.85	23058
non functional	0.82	0.71	0.76	16507
accuracy			0.82	39565
macro avg	0.82	0.80	0.81	39565
weighted avg	0.82	0.82	0.81	39565

Our model here is a bit better overall. It's a little more accurate, and overfits less


```
In [153]: sorted_idx2 = xgb_clf.feature_importances_.argsort()
plt.figure(figsize=(10,12))
plt.barh(X.columns[sorted_idx2], xgb_clf.feature_importances_[sorted_idx2])
plt.ylabel('Feature', fontsize=16)
plt.title("Feature Importance", fontsize=20)
plt.xlabel("Random Forest Feature Importance")
plt.savefig('Feature Importance2')
```



Decision Tree

```
In [150]: tree = DecisionTreeClassifier(random_state=42)
DTclf = tree.fit(X_train, y_train)

tree_training_pred = tree.predict(X_train)
tree_training_accuracy = accuracy_score(y_train, tree_training_pred)

tree_val_preds = tree.predict(X_test)
tree_val_accuracy = accuracy_score(y_test, tree_val_preds)
```

```
In [152]: print(classification_report(y_test, tree_val_preds))
```

	precision	recall	f1-score	support
functional	0.75	0.75	0.75	7694
non functional	0.65	0.66	0.65	5495
accuracy			0.71	13189
macro avg	0.70	0.70	0.70	13189
weighted avg	0.71	0.71	0.71	13189

```
In [151]: print(classification_report(y_train, tree_training_pred))
```

	precision	recall	f1-score	support
functional	1.00	1.00	1.00	23058
non functional	1.00	1.00	1.00	16507
accuracy			1.00	39565
macro avg	1.00	1.00	1.00	39565
weighted avg	1.00	1.00	1.00	39565

Conclusions

Our best model was the XGBoost. It was able to correctly guess the waterpoint status 74% of the time.

The best most important feature for identifying the state of the wells were the amount of water available to the waterpoint.

Construction year and location was also important factors.

Recommendations

Look into these water points with low amounts of access to water. Why are these not getting the water they need? If we can give these wells better access to water, we should be able to solve a large portion of the problems.

There is definitely a pattern with location and the status of the wells. We saw this in the visualizations and in our models. Try to find why this is. Is it a problem with the local regulations or something larger?

The population surrounding the wells also seems to be important. A lower population probably means that there are less regulations and maintenance. These wells are still needed though. The government needs to focus on getting these wells back online.

Future Work

We focused solely on the functioning and non functioning wells, since it was the more pressing issue. In the future we need to also better recognize which wells need maintenance, so the gap does not increase.

This data only go up to 2013. Update the data with more recent well information.

In []: